

# Amazon DynamoDB: AWS NoSQL Database for Massive Scale

It is not a relational database and does not use SQL. Instead, it is designed for key-value and document-based data models.

DynamoDB is a fully managed NoSQL database service that is highly available and automatically replicated across three Availability Zones (AZs).

## ⚡ Key Features:



Serverless architecture: No need to manage servers or infrastructure.

Handles millions of requests per second, supporting trillions of rows and hundreds of terabytes of data.

Ensures fast and consistent performance with single-digit millisecond latency for real-time applications.

Integrated with AWS IAM for fine-grained access control and secure administration.

Supports auto-scaling to adjust to workload demands and reduce costs.

## 💰 Storage Classes:

Two table types are available:

Standard: For frequently accessed data.

Infrequent Access (IA): For cost-effective storage of rarely accessed data.

DynamoDB is ideal for applications needing low-latency access at massive scale, such as mobile apps, IoT, and gaming backends.

# Understanding Data Modeling in DynamoDB (Key-Value Structure)

## 🔑 Primary Key Structure

A **Primary Key** in DynamoDB consists of:

**Partition Key** (e.g., Product ID) – Used to distribute data data across storage nodes.

**Sort Key** (optional, e.g., Type) – Used to organize related items related items under the same partition.

## 🌟 Flexible Schema per Item

Each item in a DynamoDB table can have a **different set of set of attributes**—a key advantage of NoSQL.

The **schema is defined at the item level**, not at the table level.

Different types of content (books, albums, movies) can be stored in the same table without rigid structure.

## 🧠 What This Means

The table behaves like a collection of items where:

Each item is uniquely identified by its key combination.

DynamoDB is ideal for applications needing **flexible data models**, **fast key-based access**, and **scalable storage**.

DynamoDB is a **key-value and document database**, where each item (record) is stored with a **primary key** and optional **attributes**.

In the diagram:

Product ID = 1, Type = Book ID has attributes like "Odyssey", "Homer", 1871.

Product ID = 2 is reused for two different types of items (Album and Track), grouped under the same partition.

Product ID = 3, Type = Movie ID has different attributes such as genre and director.

✓ DynamoDB is ideal for applications needing **flexible data models**, **fast key-based access**, and **scalable storage**.



# Amazon DynamoDB Accelerator (DAX): Speed Boost for NoSQL



## Application Request

Applications first send requests to **DAX**.

## Cache Check

If data is cached, DAX responds instantly.

## Database Fetch

If not, DAX fetches the data from **DynamoDB**, caches it, and returns it.

**DAX (DynamoDB Accelerator)** is a **fully managed in-memory caching layer** designed specifically for DynamoDB.

It improves read performance by up to **10x**, reducing latency from **milliseconds to microseconds**.

### Key Benefits:

**In-memory cache:** Reduces the need to repeatedly query the main DynamoDB table.

**Highly scalable and secure:** Built to handle large-scale traffic and integrated with AWS security services.

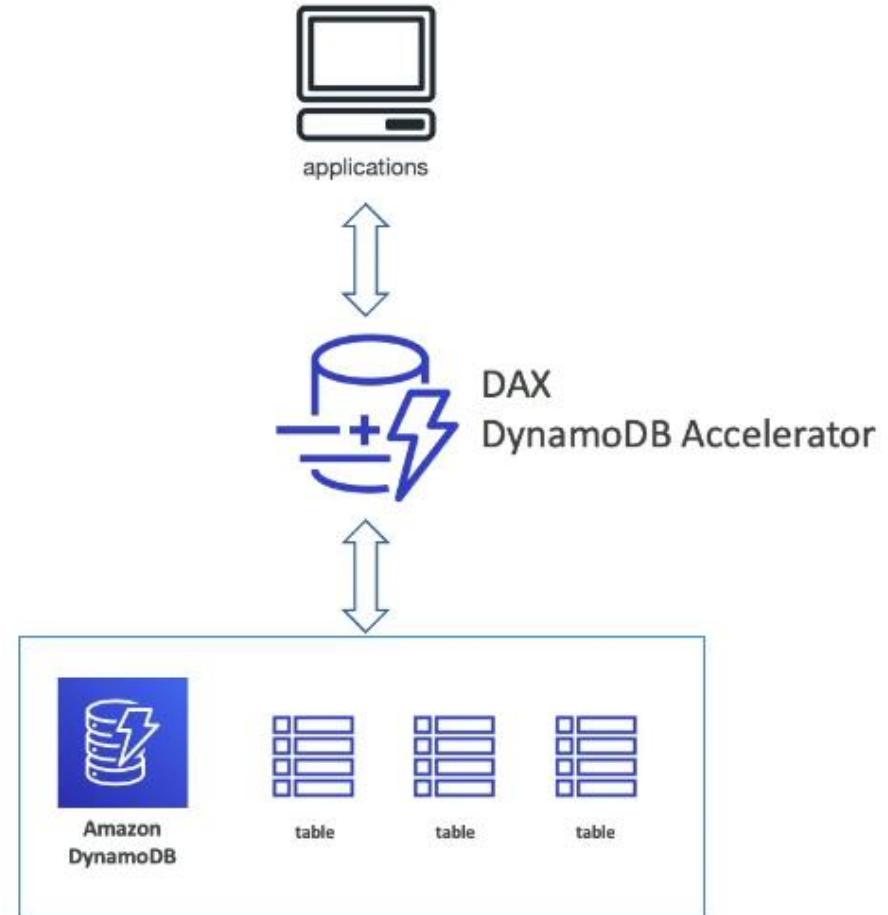
**Drop-in replacement for DynamoDB SDK:** Applications use DAX without major code changes.

### DAX vs. ElastiCache:

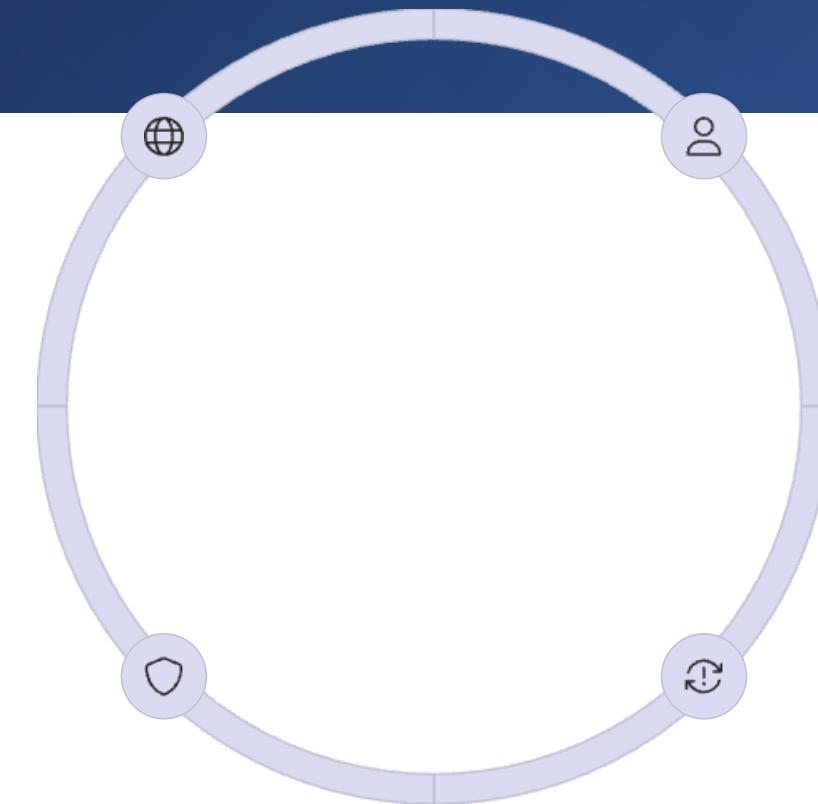
**DAX is only used with DynamoDB** and tightly integrated into its architecture.

**ElastiCache** is more general-purpose and can cache results from any type of database (e.g., RDS, Redshift, etc.).

Ideal for **read-heavy workloads** where speed is critical, such as gaming, ad tech, and mobile apps.



# DynamoDB Global Tables: Real-Time Multi-Region Access



Global Tables in DynamoDB allow data to be stored and accessed in multiple AWS regions with low latency.

This enables **read and write operations from any region** using a model called **Active-Active replication**.

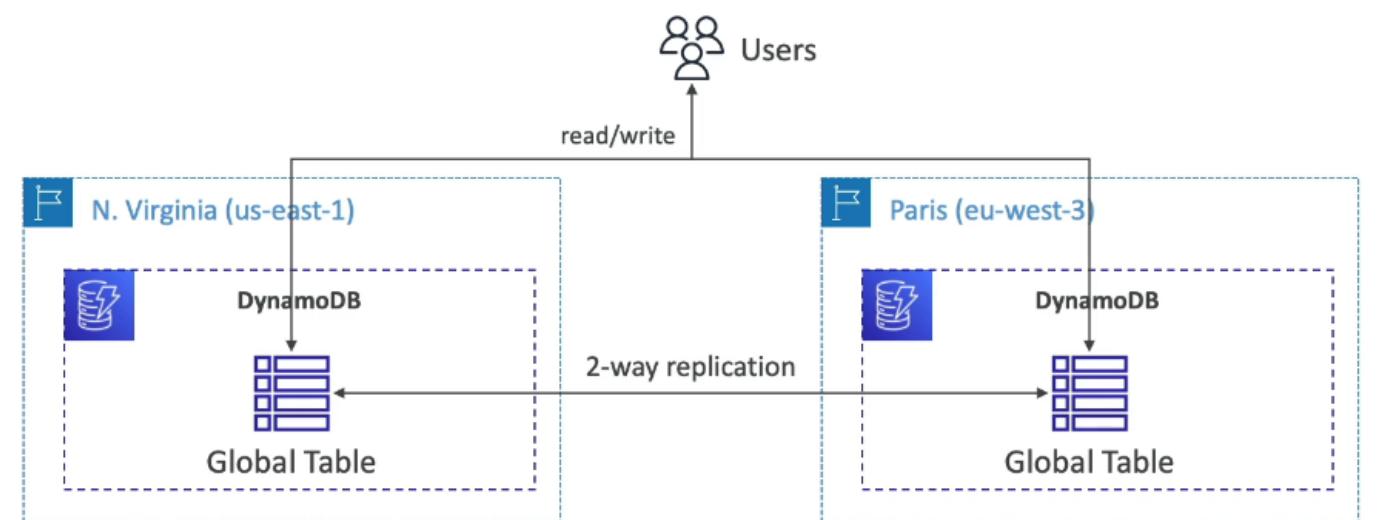
## Key Benefits:

**Low-latency access** for globally distributed applications.

**Active-Active architecture**: all regions can serve writes.

**Built-in replication** keeps data consistent across all locations.

❖ Ideal for global applications like e-commerce, gaming, and social platforms where users access data from various continents.



# What Is Docker? — A Simple Overview



## Containerization

Containers bundle everything the app needs: code, libraries, settings, system tools.



## Consistent Environments

Apps run the **same everywhere** — on laptops, servers, or cloud.



## Fast Deployment

**Fast deployment** — containers can start or stop in **seconds**.



## Easy Scaling

Easier to **scale** applications up or down as needed.

Docker is a platform that helps **developers package and run applications** inside lightweight, portable units called **containers**.

## What Are Containers?

Containers bundle everything the app needs: code, libraries, settings, system tools.

They can be **run on any system** — no matter the underlying hardware or operating system.

## Why Docker Is Powerful:

**No compatibility issues** — eliminates "it works on my machine" problems.

**Predictable performance** and **fewer bugs** due to consistent environments.

Compatible with **any language, OS, or tech stack**.

 Docker simplifies the development-to-production journey and enables rapid, reliable software delivery.



# Docker on an Operating System (OS)

A **Docker container** is a lightweight unit that includes an application and all its dependencies.

Containers run on top of a **host operating system**—for example, a Linux OS on an **EC2 instance** (virtual server).

## What the Diagram Shows:

A **server (EC2 instance)** is running **multiple containers** side by side.

Each container runs a **different application or service**, such as:

 Java-based apps (Java logo)

 Node.js apps (Node logo)

 MySQL database (MySQL logo)

All containers **share the host OS** but remain **isolated** from each other.

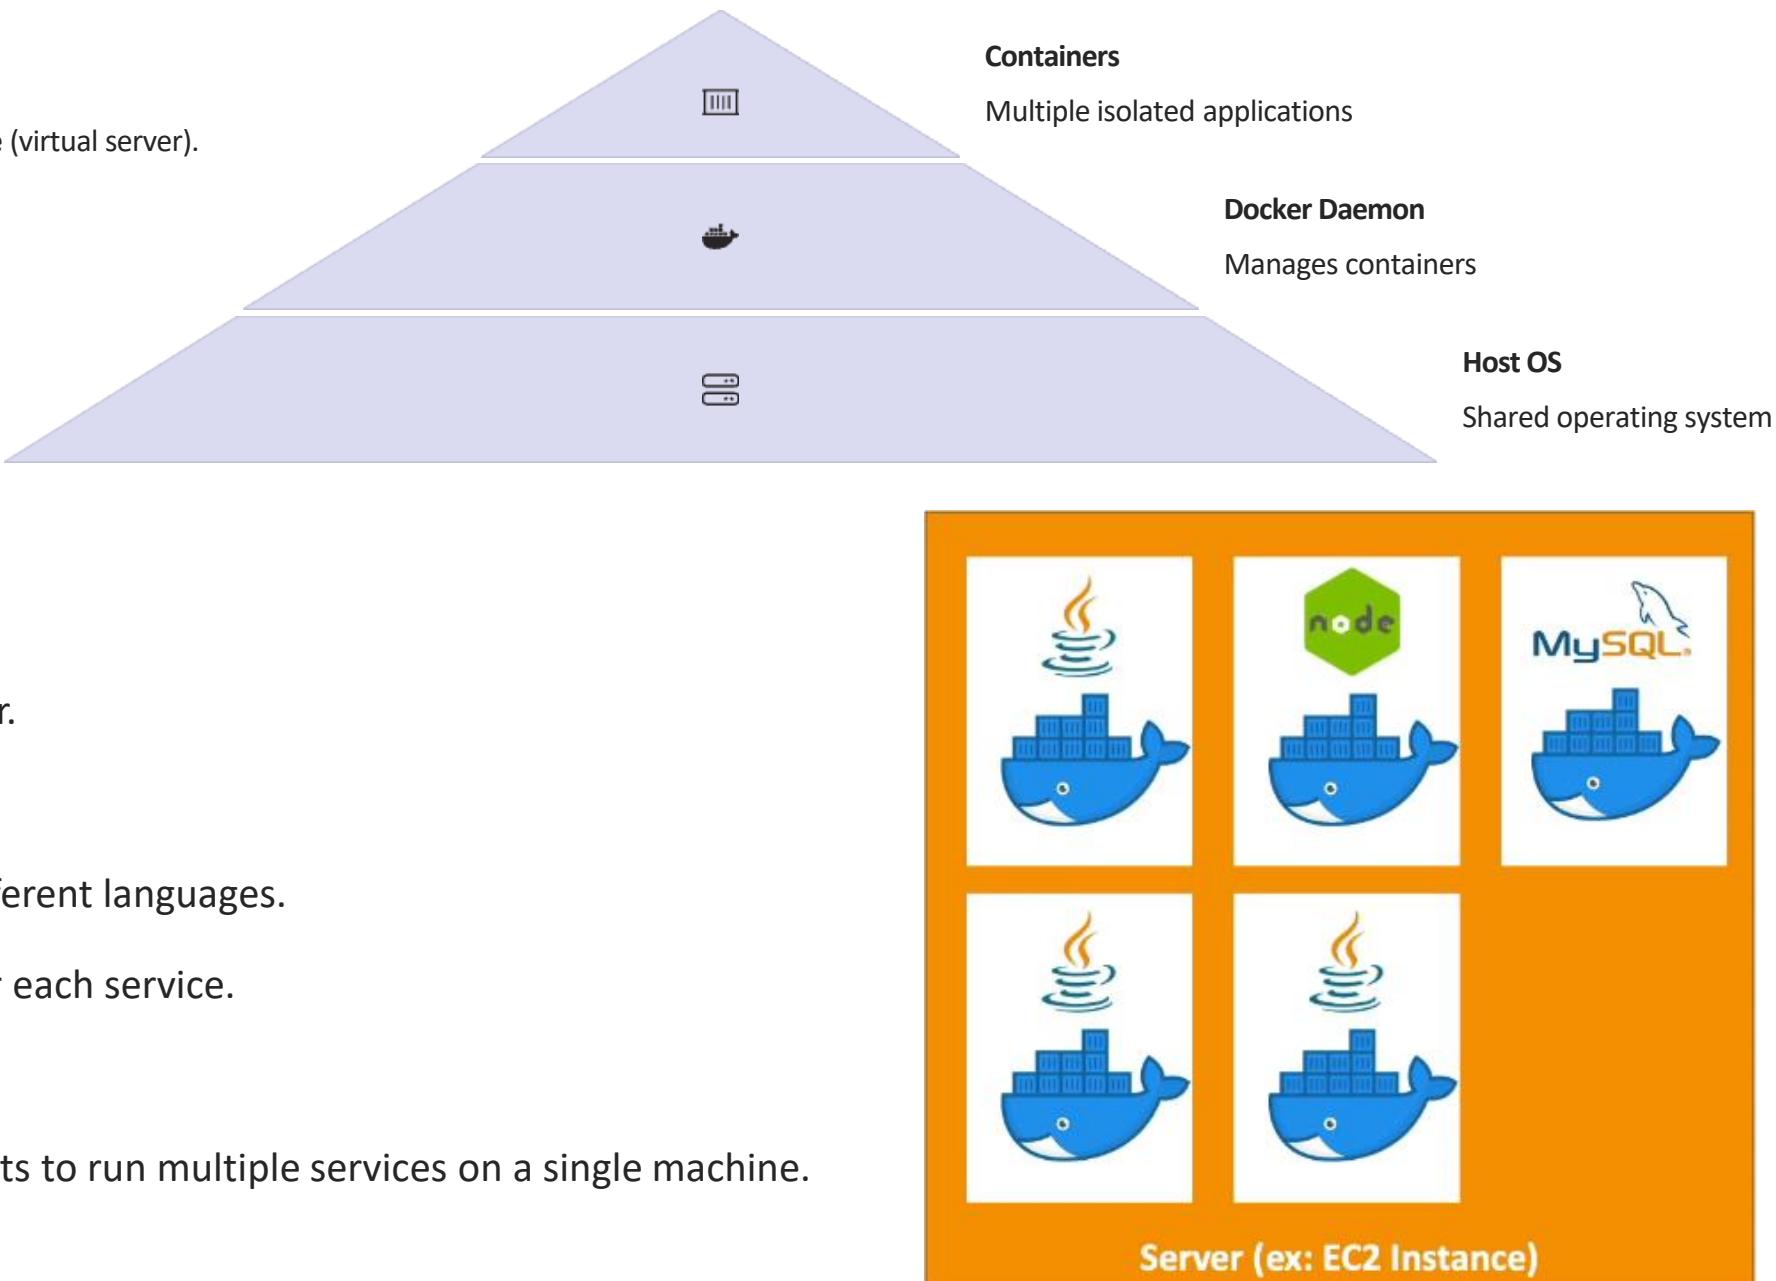
## Benefits of This Setup:

Applications are **independent** and don't conflict, even if they use different languages.

Resources are used **efficiently**—no need to spin up separate VMs for each service.

**Faster startup time, easy to scale, and easy to manage.**

 This architecture is commonly used in microservices environments to run multiple services on a single machine.



# Where Are Docker Images Stored?

## Public Repositories

Most popular: **Docker Hub**

Contains **base images** like:

- Ubuntu (Operating System)
- MySQL (Database)
- NodeJS, Java (Runtime environments)

Useful for quickly starting common applications.

Docker images are saved in **repositories**, which are locations that store, share, and manage images.

## Two Main Types of Repositories:

### 1. Public Repositories

### 2. Private Repositories

## Summary:

Docker images = packaged applications.

Repositories = cloud storage places where images live.

Public (for sharing) vs. Private (for security & control).

 When running a container, Docker pulls the image from the specified repository first.

## Private Repositories

Example: **Amazon ECR (Elastic Container Registry)**

Used in enterprise settings to securely store custom or sensitive Docker images.

Offers **access control** and **integration** with AWS services.

## Virtual Machines

### Virtual Machines vs. Docker Containers

Each VM runs **its own Guest OS** (e.g., Linux, Windows).

A **hypervisor** (like VirtualBox or VMware) sits between the Host the Host OS and VMs.

**Heavier and slower:** More memory and CPU are used due to due to multiple OS layers.

### Docker Containers

All containers **share the same Host OS**.

The **Docker Daemon** runs directly on the Host OS to manage manage containers.

**Lightweight and fast:** No need for full OS for each app.









# Clustering

- Main Clustering Ideas:
  - Use features to decide which points are most similar to other points.
  - Realize that there is no final correct **y** label to compare cluster results to.
  - We can think of clustering as an unsupervised learning process that “discovers” potential labels.



# Clustering

- Unsupervised Learning Paradigm Shift:
  - *How do we assign a new data point to a cluster?*
    - Different approaches depending on the unsupervised learning algorithm used.
    - Use features to assign most appropriate cluster.



# Clustering

- Unsupervised Learning Paradigm Shift:
  - *How do we assign a new data point to a cluster?*
    - Just as before, no way to measure if this was the “correct” assignment.