

DataWrangler

Project Overview

- ▶ In this project, we will leverage the power of data Wrangler service in AWS to prepare, clean and visualize the data.
- ▶ We will analyze the Titanic dataset which contains features related to Titanic passengers and cardiovascular disease datasets (final project).
- ▶ Here are the key learning outcomes:
 - ▶ Understand feature engineering strategies and tools.
 - ▶ Understand the fundamentals of Data Wrangler in AWS.
 - ▶ Perform one hot encoding and normalization.
 - ▶ Perform data visualization Using Data Wrangler.
 - ▶ Export a data wrangler workflow into Python Script.
 - ▶ Create a custom formula and apply it to a given column in the data.
 - ▶ Generate summary table tables in Data Wrangler.
 - ▶ Generate bias reports.



Sagemaker Data Wrangler 101

Data Wrangler 101

- Amazon SageMaker Data Wrangler accelerates the process of data preparation, exploration, cleaning, visualization and feature engineering. It makes creating Extract, Transform and Load (ETL) pipelines much easier.

The screenshot shows the Data Wrangler interface with a data preview window titled "my_first_dataflow.flow" and "Data types - Transform: titanic.csv". The preview displays a table of 26 rows from the titanic.csv dataset. The sidebar on the right lists various ETL operations:

- Custom transform
- Balance data
- Custom formula
- Encode categorical
- Feature date/time
- Feature text
- Format string
- Group by
- Handle missing
- Handle outliers
- Handle structured column
- Manage columns

The screenshot shows the Amazon SageMaker Studio interface with a "untitled1.flow" analysis window. The main area displays a histogram titled "Histogram: Height Data Visualization" with the x-axis labeled "height(binned)" and the y-axis labeled "Count of Records". The histogram bars show a peak between 160 and 170. Below the histogram is a "Data table" with columns: id, age, gender, height, and weight. The table contains 5 rows of data. To the left of the analysis window is the "RUNNING INSTANCES" panel showing an "mLm5.4xlarge" instance, and the "RUNNING APPS" panel showing "sagemaker-data-wrangler" and "untitled1.flow" both running on the same instance.

Data Wrangler 101

- ▶ SageMaker Data Wrangler is **cloud based** and **doesn't require any code!**
- ▶ Data can be imported into data Wrangler from more than one source such as **S3**, **RedShift** and **SageMaker Feature Store**.
- ▶ Data could be in **CSV, database tables and Parquet** formats.
- ▶ Data Wrangler includes over **300 data transformations** such as one hot encoding, normalization, imputation of missing data, ..etc.
- ▶ Several data visualization templates are available to generate **bar charts, line plots, histograms, scatterplots...etc.**
- ▶ Data Transformation **workflows can be exported** from data wrangler to a notebook or script so it can be automated with **SageMaker Pipelines**.
- ▶ Check out success stories from customers:
<https://aws.amazon.com/sagemaker/data-wrangler/>

What Is Feature Engineering?

- Machine Learning algorithms require training data to train.
- Feature engineering is a critical task that is performed by data scientists prior to training AI/ML models to ensure solid trained model performance.
- Feature engineering is an art of introducing new features that weren't existing before.
- Data scientists spend 80% of their time performing feature engineering.
- The remaining 20% is the easy part which includes training the model and performing hyperparameters optimization.
- As a data scientist, you may need to:
 1. Highlight important information in the data
 2. Remove/isolate unnecessary information (e.x.: outliers).
 3. Add your own expertise and domain knowledge to the alter the data.



Photo Credit: <https://pixabay.com/illustrations/network-data-memory-data-collection-4478146/>

Feature Engineering: Proper Questions To Ask?

- As a data scientist, you need to answer the following questions:

Which features should I select?

Can I add my domain knowledge to use less features?

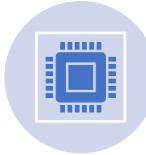
Can I come up with new features from the data I have at hand?

What should I put in the missing data locations?

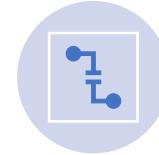
What are the capabilities of the ML model I have?



It is important to choose features that are most relevant to the problem.



Adding new features that are unnecessary will increase the computational requirements needed to train the model (curse of dimensionality).



There are many techniques that could be used to reduce the number of features (compress/encode the data) such as Principal Component Analysis (PCA) – will be covered later.

FEATURE ENGINEERING: QUIZ

- Let's take a look at this data and see what's wrong with it!

CUSTOMER ID	CUSTOMER NAME	LOCATION	CLICK ON AD?
1	Georgina	USA	Yes
2	Leila	Canada	1
3	Sarah	France	0
4	Bird		1
5	Max	Netherlands	0
6	Sarah	France	0

Feature Engineering: Solution

- Let's take a look at this data and see what's wrong with it!

CUSTOMER ID	CUSTOMER NAME	LOCATION	CLICK ON AD?
1	Georgina	USA	Yes
2	Leila	Canada	1
3	Sarah	France	0
4	Bird		1
5	Max	Netherlands	0
6	Sarah	France	0

ENTIRE COLUMN REQUIRES ENCODING

MISSING INFORMATION

REQUIRES FORMATTING

DUPLICATE ENTRY

The diagram illustrates several issues with the raw data:

- Entire Column Requires Encoding:** The "CLICK ON AD?" column contains categorical data ("Yes", "1", "0") that needs to be converted into numerical or binary values.
- Missing Information:** The "LOCATION" column for customer ID 4 is empty, which is missing information.
- Requires Formatting:** The "CLICK ON AD?" column contains mixed data types ("Yes", "1", "0").
- Duplicate Entry:** Customer IDs 3 and 6 both have the name "Sarah" and the location "France".

FEATURE ENGINEERING TECHNIQUES

Imputation

Handling Outliers

Binning

Log Transform

One-Hot Encoding

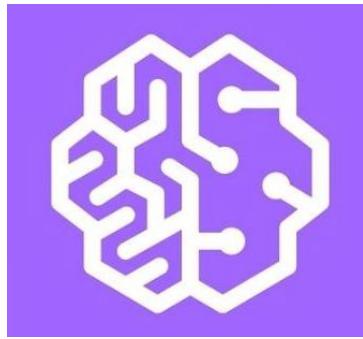
Feature Split

Scaling

Feature Engineering: Tools



JUPYTER
NOTEBOOKS



AMAZON
SAGEMAKER DATA
WRANGLER



AWS GLUE

ONE-HOT ENCODING

One-hot Encoding: Why Do We Need It?

- Can we simply replace colors with integer values?
- The machine learning model will assume that:

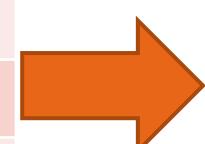
GREEN > YELLOW > RED

A diagram illustrating the problem of using raw integer encoding for categorical data like colors. On the left, a vertical list of colors is shown: RED, RED, YELLOW, GREEN, and YELLOW. On the right, a corresponding list of integers is shown: 1, 1, 2, 3, and 2. A large orange arrow points from the left list to the right list. A large red arrow, containing the word "WRONG!", points from the left list to the right list, indicating that this mapping is incorrect because it assumes a specific ordering (RED > YELLOW > GREEN) that does not reflect the true meaning of the categories.

COLOR	ENCODED COLOR
RED	1
RED	1
YELLOW	2
GREEN	3
YELLOW	2

One-hot Encoding

- One hot encoding is widely used in machine learning.
- It works by converting values such as “color” into columns with 1’s and 0’s in them.
- Since machine learning models deal with numbers, we perform one hot encoding to convert from categorical data into numerical.
- If you have N categories, you will need N-1 binary columns to represent them.



COLOR	RED	YELLOW	GREEN
RED	1	0	0
RED	1	0	0
YELLOW	0	1	0
GREEN	0	0	1
YELLOW	0	1	0

ONE-HOT ENCODING: ORDINAL Vs. Nominal

- The difference between nominal and ordinal data is as follows:
 - In ordinal data, order is important.
 - In nominal data, order is not important.

NOMINAL

Order of colors doesn't mean anything!

COLOR
RED
RED
YELLOW
GREEN
YELLOW

ORDINAL

Order is important!



- 1 star means poor quality course
- 5 star means great quality course

Photo Credit: <https://pixabay.com/vectors/rating-stars-system-evaluation-153125/>

Feature Scaling

- Feature Scaling is an important step to take prior to training of machine learning models to ensure that features are within the same scale.
- Example: interest rate and employment score are at a different scale. This will result in one feature dominating the other feature.
- Scikit Learn offers several tools to perform feature scaling.

RAW ORIGINAL DATASET

	Interest Rates	Employment	S&P 500 Price
0	1.943859	55.413571	2206.680582
1	2.258229	59.546305	2486.474488
2	2.215863	57.414687	2405.868337
3	1.977960	49.908353	2140.434475
4	2.437723	52.035492	2411.275663
5	2.143637	56.060598	2187.344909
6	2.148647	51.513208	2263.049249
7	2.176184	53.475909	2281.496374
8	2.125352	63.668422	2355.163011
9	2.225682	56.993396	2326.330337
10	1.814688	55.361780	2078.553895
11	2.281897	58.484752	2337.504507
12	2.426738	55.709328	2485.774097

QUICK STATS!

	Interest Rates	Employment	S&P 500 Price
count	1000.00	1000.00	1000.00
mean	2.20	56.25	2320.00
std	0.24	4.86	193.85
min	1.50	40.00	1800.00
25%	2.04	53.03	2190.45
50%	2.20	56.16	2312.44
75%	2.36	59.42	2455.76
max	3.00	70.00	3000.00

Normalization

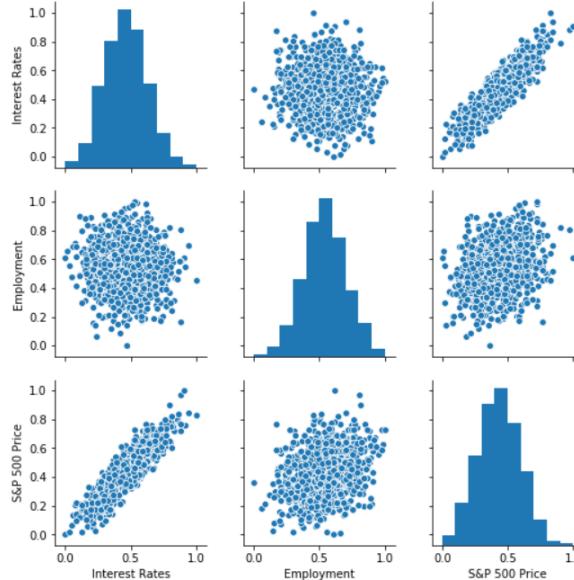
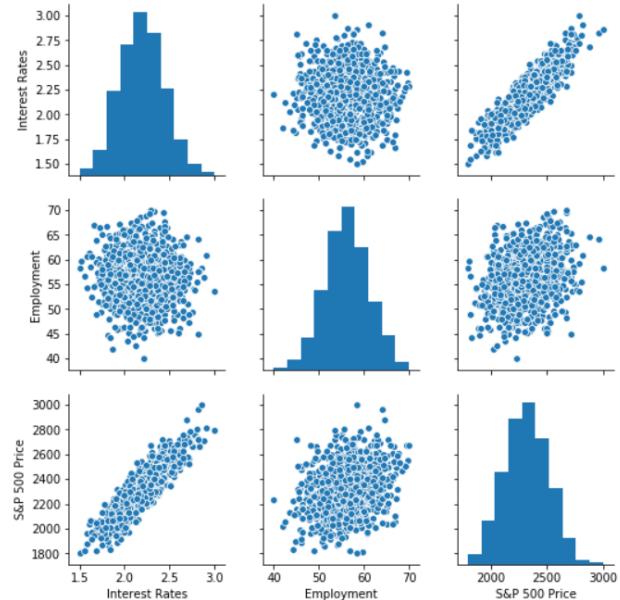
- Normalization is conducted to make feature values range from 0 to 1.

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
stock_df = scaler.fit_transform(stock_df)
```

Normalization

- Normalization is conducted to make feature values range from 0 to 1.

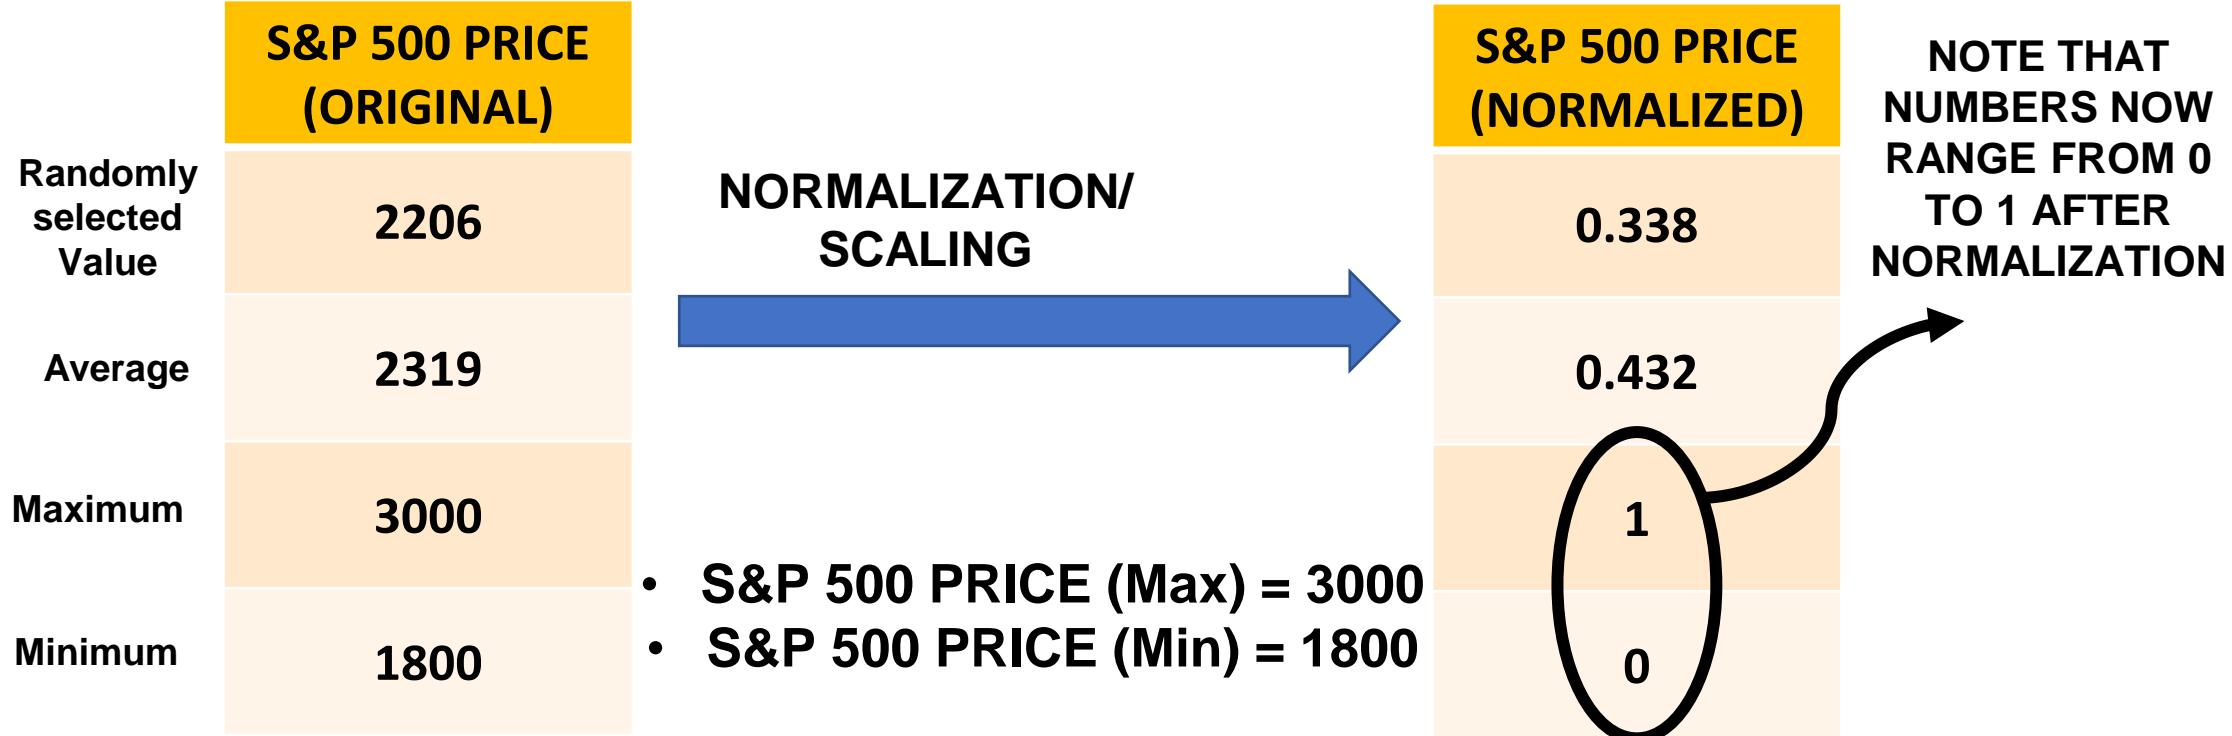


	Interest Rates	Employment	S&P 500 Price
count	1000.00	1000.00	1000.00
mean	2.20	56.25	2320.00
std	0.24	4.86	193.85
min	1.50	40.00	1800.00
25%	2.04	53.03	2190.45
50%	2.20	56.16	2312.44
75%	2.36	59.42	2455.76
max	3.00	70.00	3000.00

	Interest Rates	Employment	S&P 500 Price
count	1000.00	1000.00	1000.00
mean	0.46	0.54	0.43
std	0.16	0.16	0.16
min	0.00	0.00	0.00
25%	0.36	0.43	0.33
50%	0.47	0.54	0.43
75%	0.57	0.65	0.55
max	1.00	1.00	1.00

Normalization

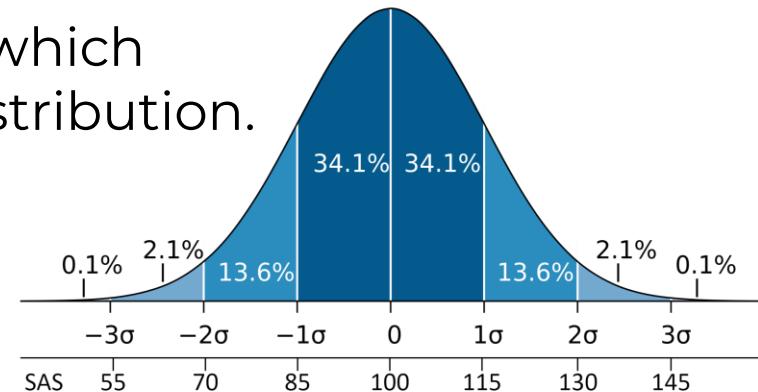
$$x' = \frac{x - \min(x)}{\max(x) - \min(x)} = \frac{2206 - 1800}{3000 - 1800} = 0.338$$



Standardization

- Standardization is conducted to transform the data to have a mean of zero and standard deviation of 1.
- Standardization is also known as Z-score normalization in which properties will have the behaviour of a standard normal distribution.

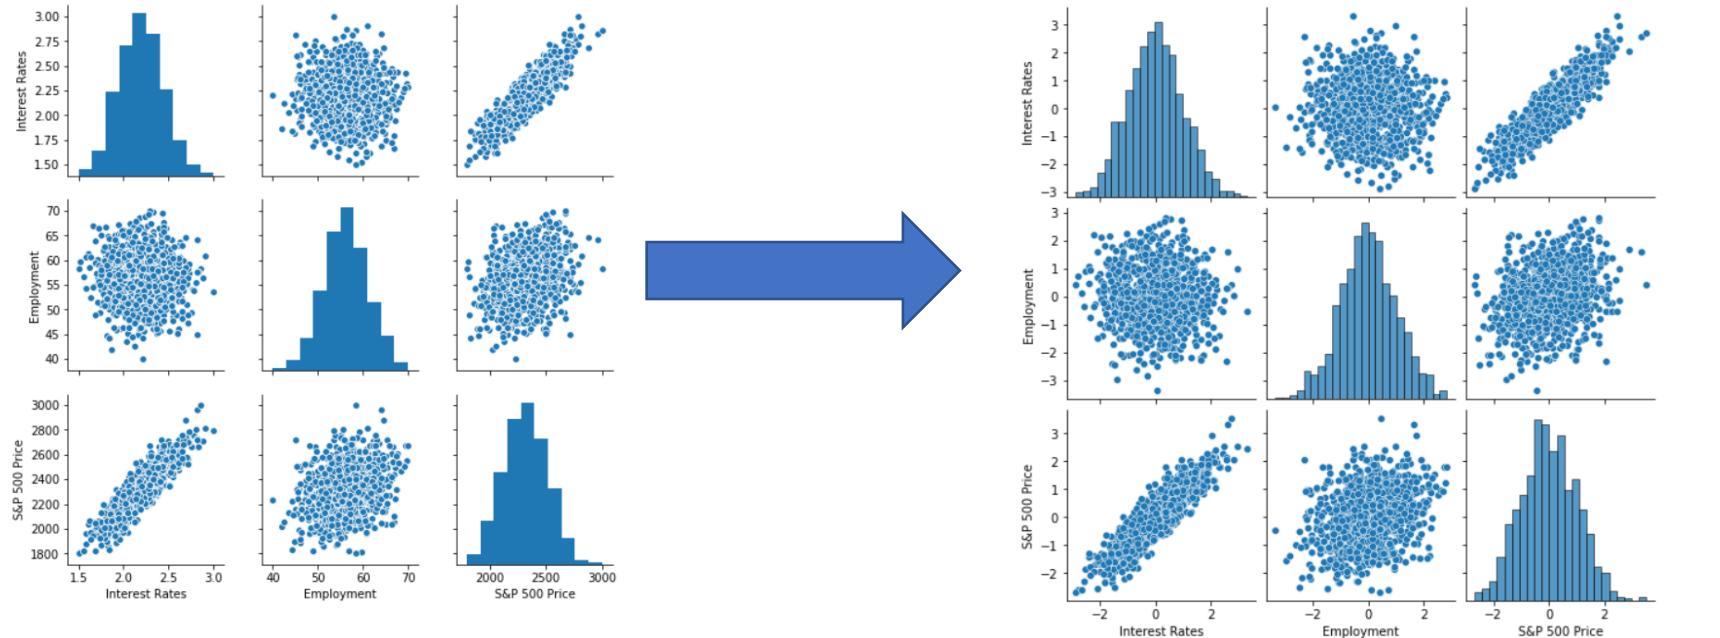
$$z = \frac{x - \bar{x}}{\sigma}$$



```
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
stock_df = scaler.fit_transform(stock_df)
```

Standardization

- Standardization transforms data to have a mean of zero and standard deviation of 1.



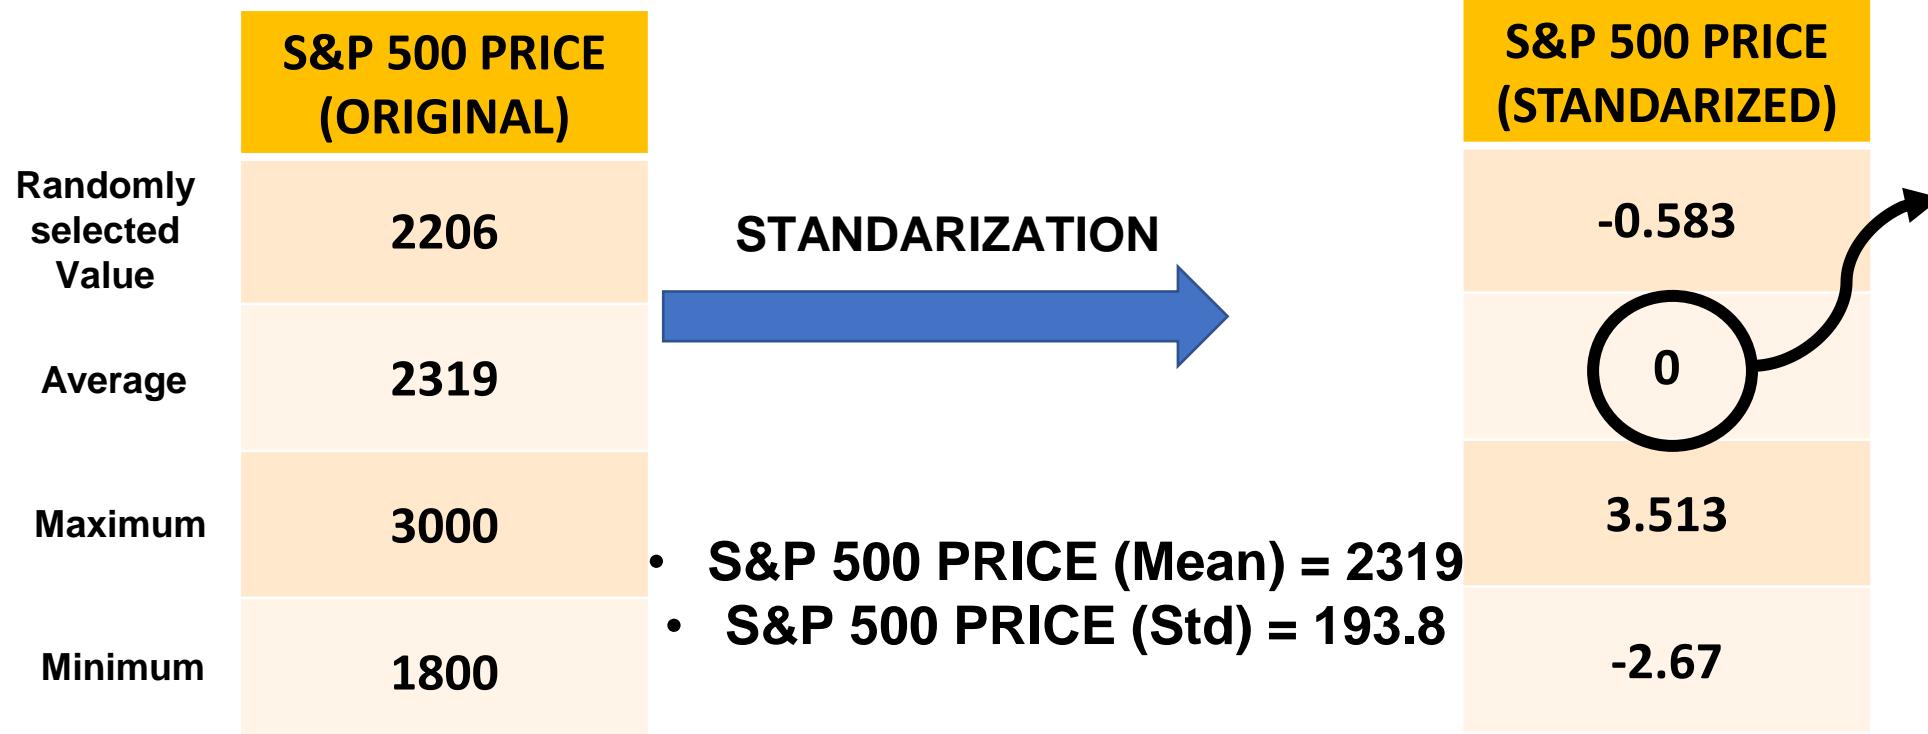
	Interest Rates	Employment	S&P 500 Price
count	1000.00	1000.00	1000.00
mean	2.20	56.25	2320.00
std	0.24	4.86	193.85
min	1.50	40.00	1800.00
25%	2.04	53.03	2190.45
50%	2.20	56.16	2312.44
75%	2.36	59.42	2455.76
max	3.00	70.00	3000.00

The figure displays a 3x3 grid of plots illustrating the standardization process for three variables: Interest Rates, Employment, and S&P 500 Price. The left column shows histograms of the original data. The middle column shows scatter plots of the data pairs. The right column shows histograms of the standardized data. A large blue arrow points from the original data to the standardized data.

	Interest Rates	Employment	S&P 500 Price
count	1000.00	1000.00	1000.00
mean	0.00	0.00	-0.00
std	1.00	1.00	1.00
min	-2.88	-3.34	-2.68
25%	-0.66	-0.66	-0.67
50%	0.01	-0.02	-0.04
75%	0.68	0.65	0.70
max	3.33	2.83	3.51

Standardization

$$z = \frac{x - \bar{x}}{\sigma} = \frac{2206 - 2319}{193.8} = -0.583$$



NOTE THAT AFTER STANDARDIZATION THE AVERAGE IS SET TO ZERO

Always Remember!

“A normalized dataset will always range from 0 to 1”

“A standardized dataset will always have a mean of 0 and standard deviation of 1, but can have any upper and lower values”

WHEN SHOULD I PERFORM STANDARDIZATION VS. NORMALIZATION?

Scaling (standardization or normalization) is required when we use any machine learning algorithm that require **gradient calculation**.

Examples of machine learning algorithms that require gradient calculations are: linear/logistic regression and artificial neural networks

Having different scales for each feature will result in a different step size which in turn jeopardizes the process of reaching a minimum point.

Scaling is not required for distance-based and tree-based algorithms such as K-Means Clustering, Support Vector Machines and K Nearest Neighbors, decision trees, random forest, and XG-Boost.

STANDARDIZATION Vs. Normalization?

Generally speaking,
there is no right or
wrong answer!

In case of neural
networks, normalization
is preferred since we
don't assume any data
distribution.

Standardization is
preferred when data
follows gaussian
distribution

Standardization is
preferred over
normalization when
there are a lot of
outliers.

Thanks so much
to

Professor Ryan Ahmed

